

Operating Systems

Unit 4

File-System Interface

* File → a logical storage unit

(i) text file : a sequence of characters organized into lines and pages

(ii) source file : a sequence of functions, organized as declarations and executable statements

(iii) executable file:- a series of code sections that the loader can bring into the memory and execute.

File Attributes

→ name → protection

→ identifier → time, date & user identification

→ type

→ location

→ size

File Operations

(i) creation

(ii) write file

(iii) read file

(iv) repositioning aka seek

(v) deleting a file

(vi) truncating a file

* Opening Files

open file table : has info about the files that are open

no need to search

file pointer points to last read / location

File open count : when files are closed, decrement counter

allows removal data from the open-file table

when the last process closes it

* Open File locking

→ to prevent other processes from gaining access to it

Shared lock - reader lock - several processes can acquire the lock concurrently

Exclusive lock - only one process can acquire lock at a time, writer's lock

Mandatory lock - once an exclusive lock is acquired, OS prevents any other process from accessing it

Advisory locks - upto developer to find lock status & decide what to do

File Access Methods

- (i) Sequential access : info is processed in order, one record after the other
 read-next()
 write next()

(ii) Direct access

- file is made up of fixed-length logical records that allow programs to read and write records rapidly in no particular order
- based on disk model, since disks allow random access to any file block

(iii) with an index

- maintain index for each file (like a book)
- contains pointers to different blocks

* Directory and Disk Structure

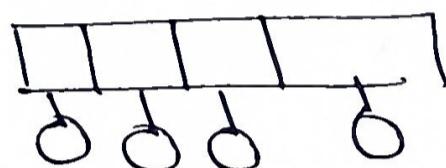
- disk divided into partitions
- partitions can be raw - w/o a file system or formatted - w/ a file system
- an entity w/ a file system called a volume
- each volume has the device directory and table of contents
-

* Directory operations

- (i) search
- (ii) create
- (iii) read
- (iv) delete
- (v) list
- (vi) rename
- (vii) traverse

* Logical Structures of Directories

(1) Single - Level



→ single directory for all users

→ naming problem

→ grouping problem

→ difficult to search

(2) Two - Level Directory

→ separate directory for each user.

→ each has their own UFD - User File Directory

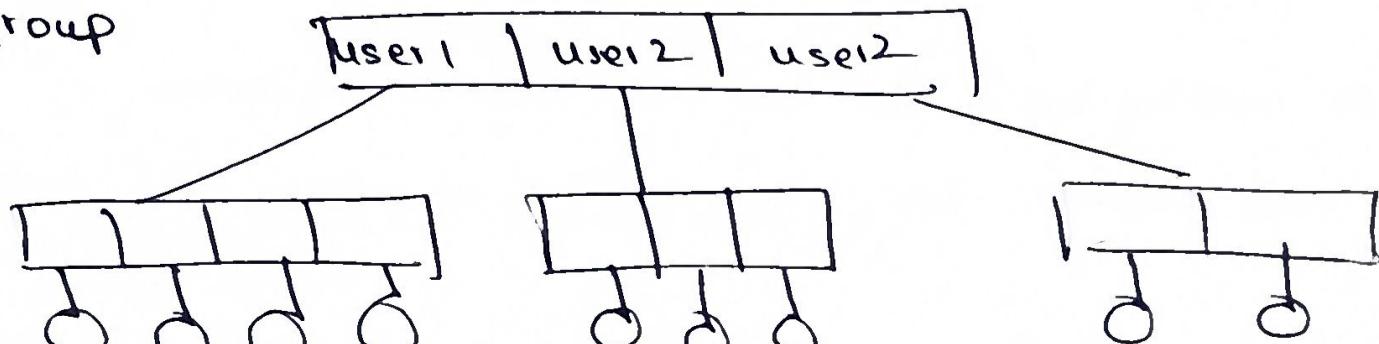
→ search for UFD in Master File Directory - MFD

→ accessed w/ path name

→ can have same file name for diff. users

→ easy to search

→ cannot group



(3) Tree - Structured Directory

- users can create their own subdirectories & organize their files accordingly.
- one bit defines the entry:
 - file = 0
 - subdirectory = 1
- use file path name
 - absolute - from root
 - relative - from cwd
- efficient searching
- grouping capacity

(4) Acyclic Graph Directories

- allows for the sharing of subdirectories
 - changes made by 1 person visible to other immediately
 - implemented using links
 - ↳ a pointer to another file / subdirectory
 - resolve the link → use path name to locate real file
- | | | |
|----------------------------------|-----|---|
| <u>Hard links</u> | vs. | <u>Soft links</u> |
| ↓ | | ↓ |
| direct pointers to original node | | a file that has the information to point to another file / inode. |

(5) General Graph Directory

- may have cycles
- would lead to infinite loops

How to guarantee no cycles?

- allow links to only files not subdirectories
- garbage collection - determine when the last reference has been deleted and the disk space can be reallocated
 - involves traversing entire file system, marking everything that can be accessed
- use a cycle detection algorithm

* File System Mounting

- a FS ^{is} must be mounted before it is accessible to the processor on the system
- need name of device & mount point

* Protection

- access control lists
 - specifies user names & types of access allowed
 - when a user requests access, OS checks

3 classifications : (i) owner - who made file
(ii) group - users sharing file
(iii) universe - all other users

* File Sharing

(i) Multiple Users

implement sharing w/ classification, owner, user & group
 user ids
 group ids

(ii) Remote File Systems

- use ftp (for anonymous & auth access)
- use a distributed file system
- via the web

client-server model

client seeks files from server

NFS, CFS

Failure modes

- corruption of dis. structures - metadata
- RFS have even more failure modes
 - network failure
 - server failure
- maintain state information
- stateless DFS - assumes that a client request wouldn't have happened unless the file was already open.
 - easy but insecure.

* File-System Structure

- File systems stored on disks because:
 - (i) disks can be rewritten in place
 - (ii) disk can directly access any block of info it contains - can access file either sequentially or randomly
- I/O transfers between memory & disk happen in blocks
- A basic file system issues generic read & write commands
- FCB - contains information about the file, including ownership, permissions & location of file contents
- use a layered structure to avoid duplication
 - Layers (i) basic file system - read, write operations
 - (ii) device driver - manage I/O devices
 - (iii) file organization module - understand files, logical address & physical blocks

* File-System Implementation

- (i) boot control block - has info needed to boot an OS
- (ii) volume control block - has info about the no. of blocks in the partition, size of blocks
- (iii) directory structure
- (iv) FCB per file - file details

* Partitions and Mounting

- a disk can be sliced into multiple partitions
- raw / cooked partitions
 - ↓
 - has no filesystem
- root partition - contains the OS, mounted at boot time.
- boot loader - loads kernel & starts executing it

* Virtual File System

- provides an object oriented way of implementing file systems
- separates file-system generic operations from their implementation
- provides a mechanism to uniquely represent a file throughout a network - vnode
- In Linux, the 4 main object types for VFS are:
 - (i) inode
 - (ii) file object
 - (iii) superblock object
 - (iv) directory object

* Directory Implementation

- (i) Linear list - time consuming, simple
- (ii) Hash Table - linear list w/ hash data structure
- decreases search time, collisions

* Allocation methods

- (i) Contiguous Allocation - each file occupies contiguous blocks
- simple
 - need to find space
 - need to know filesize
 - external fragmentation
 - compaction
- (ii) Extent Based Systems
- modified contiguous allocation scheme
 - allocates in contiguous blocks of disks
 - a file has 1 or more extents
- (iii) Linked Allocation - each file a linked list of blocks
- file ends at null pointer
 - no compaction, external frag.
 - not reliable
 - takes time to locate block
- (iv) FAT variation - beginning of volume has table indexed by block no
- (v) Indexed Allocation- each file has its own index blocks
- random access
 - need index table
 - overhead of index block

(vi) Combined Scheme

- say there are 15 ptu
- first 12 = direct blocks
- next 3 = indirect blocks
 - (i) ^{single indirect} points to has addresses of blocks w/data
 - (ii) double indirect - addresses of blocks that contain addresses of blocks
 - (iii) triple ~~address~~ indirect block

* Free Space Management

- keep track of free disk space
- maintain free-space list

① Bit Vector

- each block is 0 or 1
- 1 \Rightarrow free
- 0 = allocated

- simple to implement
- easy to find first n consecutive free blocks

② linked list

- maintain a linked list of free blocks
- cannot get continuous space easily
- no waste of space
- no need to traverse entire list

③ Grouping

→ modify linked list to store address of the next n-1 free blocks

④ Counting

→ keep address of first free block & count of following free blocks

→ free space list has entries containing addresses and counts

⑤ Space Maps

→ used in Oracle's ZFS filesystem

→ create metaslabs to divide space on disk into manageable size

→ each metaslab has an associated space map.

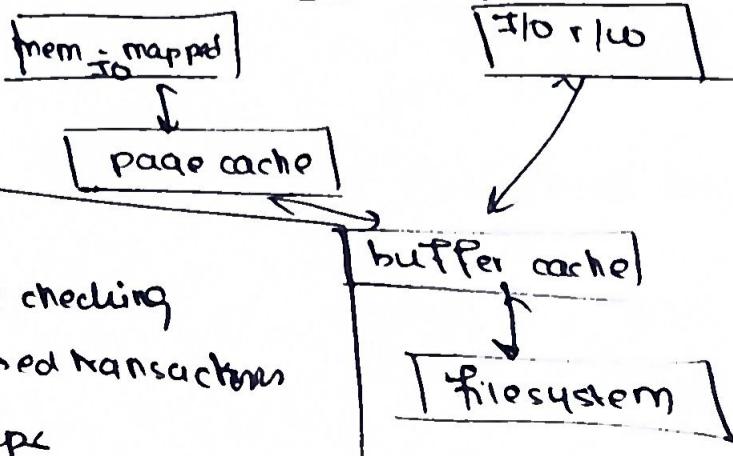
→ use the counting algo to store info about free blocks

→ use a log

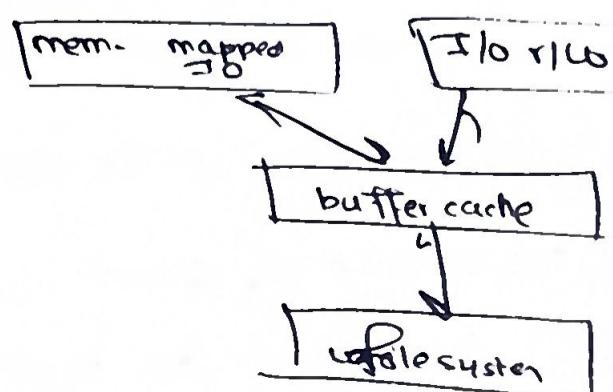
* I/O w/ a Unified Buffer Cache

→ a UBF uses the same page cache to cache both memory mapped pages and ordinary file system I/O to avoid double caching.

w/o UBF



w/ UBF

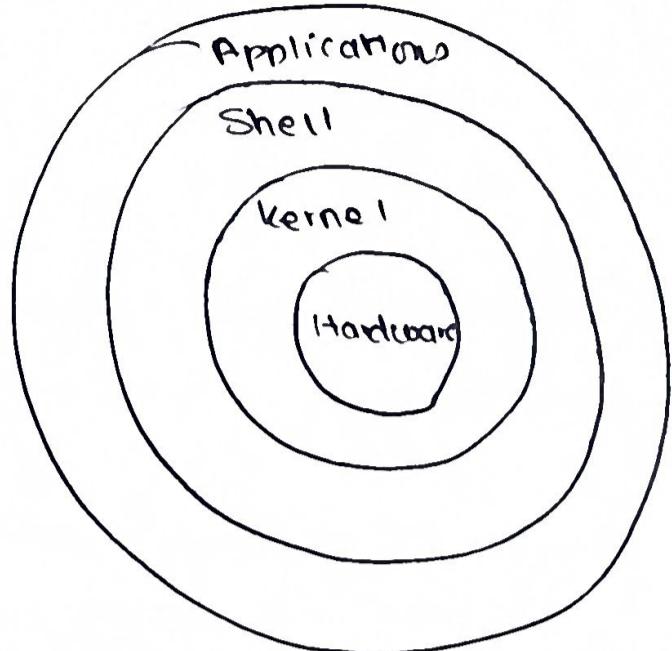


Recovery

- (i) consistency checking
- (ii) log based transactions
- (iii) backups

* Linux Architecture

- (i) multiuser
- (ii) preemptively multitasking
- (iii) adheres to traditional UNIX semantics



Kernel → responsible for maintaining / managing system resources & providing services to applications

- monolithic kernel - all core OS functionality is contained in a single executable file
- no unnecessary context switches
- highly modular and configurable
- all kernel code executes in the processor's privileged mode with full access to all the physical resources of the computer. (Kernel mode)
- user mode - has access to only a controlled subset of system's resources

Kernel Modules

module mgmt - load modules into memory com.m. w/ rest of kernel

module loader & unloader - user mode utilities

driver-reg system - allows modules to tell kernel that a new driver is available