

Unit 4

Concurrency Control

* Concurrency control: a procedure to manage multiple transactions without conflicting with each other

- helps enforce isolation through mutual exclusion
- helps w/ consistency preservation of transaction
- resolve read - write and write - write conflict issues
- enforces serializability

* Types of concurrency control protocols

- (i) locking-based protocols
- (ii) timestamp-based protocols
- (iii) validation-based protocols

* Locks

- an operation which secures permission to Read or Write a data item for transaction
- each lock is associated w/ a data item

$\text{LOCK}(x)$: Data item x locked by requesting transaction

$\text{UNLOCK}(x)$: Data item x is made available to all other transactions

* Types of locks

(i) Binary locks:

→ has only 2 states, locked or unlocked
(1 and 0)

$\text{lock}(X) = 1 \Rightarrow$ cannot be accessed

~~$\text{lock}(X) = 0 \Rightarrow$ can be accessed~~

Rules to follow (Enforced by lock manager)

1. Must issue $\text{lock}(X)$ before read / write operations
2. must issue $\text{unlock}(X)$ after all read / write operations are complete
3. cannot issue a ~~lock & unlock~~ ^{unlock} if it unless it already holds the lock on X.

Algorithm

||||| LOCK

B: if $\text{lock}(x) = 0$
then $\text{lock}(x) \leftarrow 1$

else

begin
wait until $\text{lock}(x) = 0$
go to B

end

UNLOCK

$\text{lock}(x) \leftarrow 0$

if any transactions are waiting, wake them up

Disadvantage

- at most one transaction can hold a lock on a given item
- not practically useful

(ii) Shared | Exclusive locks

3

Shared Locks - when a read operation is to be done
(Read)

Exclusive locks - when a read/ write operation is to be done.
(Read (write))

→ With read locks, other transactions are also allowed to read the item

→ With write locks, a single transaction exclusively holds the lock on the item

request →

	S	X
S	Yes	No
X	No	No

Rules
mm

1. Issues the `read-lock(X)` or `write-lock(X)` before read/write operations.
2. Issue the `unlock(X)` after all read & write operations are completed.
3. cannot issue a `read-lock(X)` if it already holds a read/write lock
(exception - downgrade from write to read)
4. cannot issue a `write-lock(X)` if it already holds a read lock or a write lock
(exception - upgrade lock from read to write)

* Lock Manager

- stores the identity of a transaction locking a data item, lock mode and the pointer to the next data item
- implemented using a linked list

Pseudocode

read-lock(x)

B: if $\text{lock}(x) = \text{unlocked}$

then begin

$\text{lock}(x) \leftarrow \text{read-locked}$

$\text{no-of-reads}(x) \leftarrow 1$

end

else if $\text{lock}(x) = \text{read-locked}$

then

$\text{no-of-reads}(x) \leftarrow \text{no-of-reads}(x) + 1$

else

begin

wait until $\text{lock}(x) = \text{'unlocked'}$

wake up transaction

go to B

end

white-lock(x):

Pseudocode:

```
B: if lock(x) = unlocked
    then lock(x) ← write-locked
else
begin
    wait until lock(x) = unlocked
    goto B
end,
```

UNLOCK(x)

```
if lock(x) = write-locked
then
begin
    wakeup one of the waiting transactions
end
else, if lock(x) = read-locked
then
begin
    no-of-reads(x) ← no-of-reads(x)-1
    if no-of-reads(x) = 0
    then
begin
        lock(x) = unlocked
        wake up transaction
end
end
```

* Conversion of Locks

- A transaction that already holds a lock on an item X is allowed under certain conditions to convert the lock from one locked state to another.
- Upgrading a lock: can upgrade if T is the only transaction holding a read lock, otherwise must wait

* 2PL

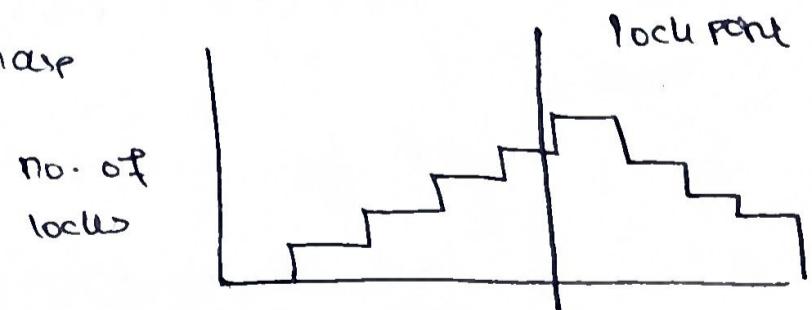
- all locking operations precede the first unlock operation
 - ensures serializability

Expanding or growing phase - new locks can be acquired, none can be released.

Shrinking - existing locks can be released, no new locks can be acquired.

Upgradation can happen during the expanding phase

downgrading during shrinking phase



Enforcing Serializability

X must remain locked by T until all r/w ops are done

Another transaction seeking to access X may be forced to wait

Also if an ~~been~~ item Y is locked earlier than it is needed, another transaction seeking access to Y is forced to wait, even though T is not using Y yet

∴ 2PL suffers from deadlock & cascading rollback

+ Types of QPL

7

1. Basic QPL - Transaction locks data items incrementally.
results in deadlock

2. Strict QPL

- does not release any write locks until after a commit
 - not deadlock-free, but no cascading rollback
 - possible to enforce recoverability

3 Rigorous QPL

- does not release any of its locks until after commit/
abort
↑
rw
 - more restrictive than strict QPL, but easier to implement
 - no cascading rollback or deadlock

4. Conservative QPL

- requires a transaction to lock all the items it access before the transaction begins execution, by predeclaring its read & write set
 - Transaction is in shrinking phase after it starts
 - is deadlock free
 - is difficult to use because of difficulty in predeclaring the read and write set.

Recovery Subsystems

- * Recovery Process : restore database to the most ^{recent} "consistent" state
- * Recovery Strategies:
 - (i) restore backed-up copy of database from archival storage
 - (ii) identify any changes that may cause an inconsistency and apply undo or redo
- * Recovery is done w/ log files (refer structure in transactions notes)

The old value of an item = Before Image = BFIIM → for UNDO
new value = After Image = AFIIM → for REDO

Caching

→ db items being read & written in disk blocks

→ buffer has modified bit, pin/ unpin bit



dirty bit, D

If buffer has been modified

set to 1 if it cannot be written back into the disk yet

(usually waiting for a commit)

* Strategies for Flushing the Buffer

(i) In-place updating

- write buffer back to orig location
- overwrites old value
- recovery requires a log

(ii) Shadowing

- write updated buffer to a diff. location
- old value = BEIM
- new value = AFIM

* Write-Ahead Log (WAL)

- (i) undo-type log entries - include old values - BEIM
- (ii) redo-type log entries - include the new values - AFIM
- When a data item is updated, entries are made in the appropriate log entry.
- With WAL, info needed for recovery must be first written onto the log file on disk, before changes are made to the database on disk.

WAL Rules

- (i) BEIM cannot be overwritten by AFIM, until all UNDO log records have been force-written onto the disk
- (ii) commit of a transaction cannot be completed until all the REDO & UNDO log records for that transaction have been force-written to disk (STABLE STORE)

* Techniques used to move a page from cache to memory

- (i) no-steal approach: can't be written to disk before commit
- (ii) steal approach: can be written to disk before commit
- (iii) Force approach: all pages updated by a transaction are immediately written to disk
- (iv) no-force approach: all pages updated after commit

* Check pointing

checkpoint - a type of entry in the log

Due to system failure, and when the system restarts

If the order of operations is

- (i) Transaction → commit → Checkpoint : no action to be taken
- (ii) Transaction → Checkpoint → commit : redo
- (iii) Transaction → Checkpoint
or
Checkpoint → Transaction : undo

see also on
Pg. 16

* Recovery Policies for Non-Catastrophic Transaction

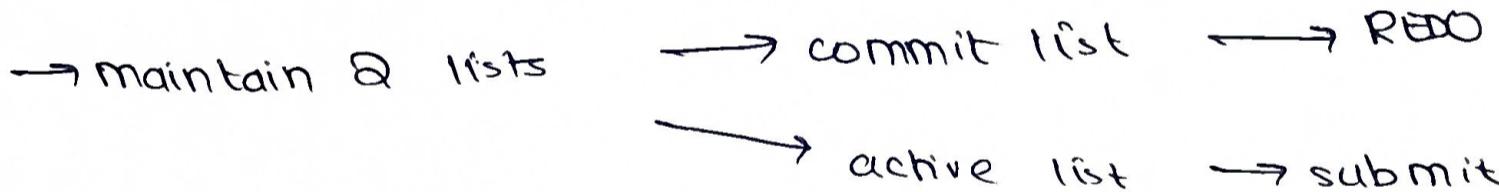
Failures

① Deferred Update

mmm mmm

- do not physically update until commit
- till commit, updates are recorded in the log file
- if failure, no changes in db so no UNDO
- after commit - REDO operations → change AFIM value
called the NO-UNDO/REDO algorithm

For concurrent users:

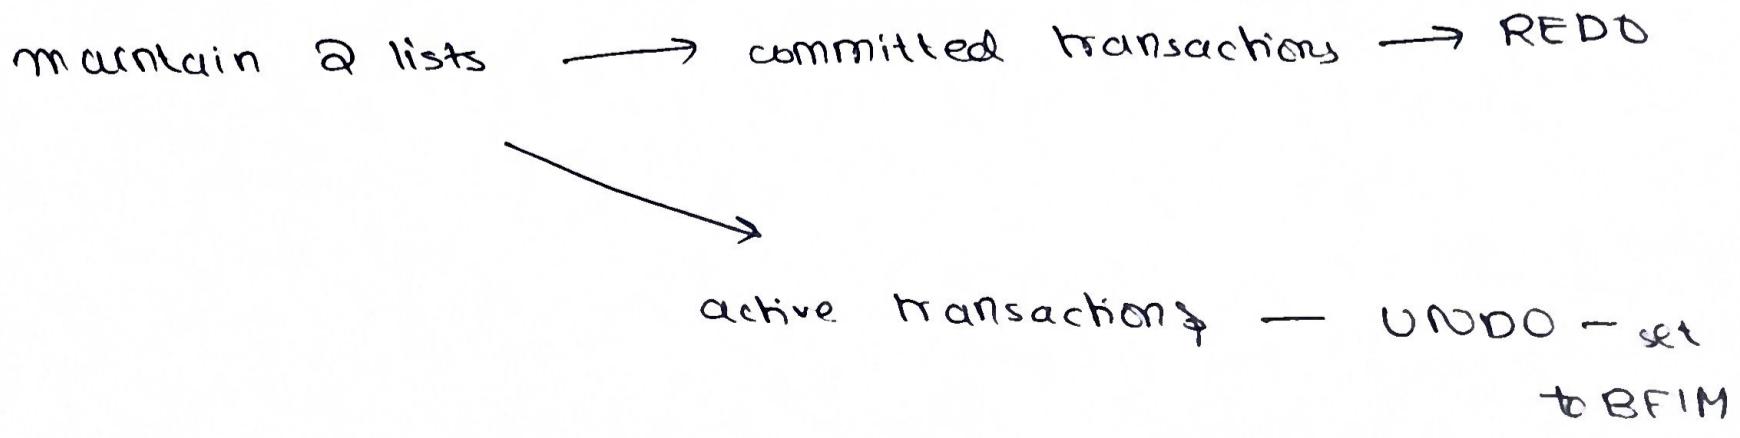


again, since they were effectively

② Immediate Update

mmm mmm

- updated for every transaction, even before commit
- in case of failure - rollback



do in reverse order of coq entry

QPL and Recovery Mechanism Problems

① Consider the 2 transactions T_1 and T_2

$T_1 : r_1(y); r_1(x); w_1(x),$

$T_2 : r_2(x); w_2(x);$

Write a serializable schedule that follows basic QPL.

T_1	T_2
$S_LOCK(y)$	
$r_1(y)$	
	$X_LOCK(x)$
	$r_2(x)$
	$w_2(x)$
	$UNLOCK(x)$
$X_LOCK(x)$	
$r_1(x)$	
$w_1(x)$	
$UNLOCK(x)$	
$UNLOCK(y)$	

② Implement strict QPL on the following transactions

$T_1 : r_1(y); r_1(x); w_1(x); \text{commit}$

$T_3 : r_3(x); r_3(y); w_3(x); \text{commit}$

T_1	T_2
$S_LOCK(y)$	
$r_1(y)$	
	$LOCK(x)$
	$r_3(x)$
	$S_LOCK(y)$
	$r_3(y)$
	$w_3(x)$
	$commit$
	$unlock(x)$
	$unlock(y)$

$\text{lock}(x)$
 $\text{unlock}(x)$
 $r_1(x)$
 $w_1(x)$
 commit
 $\text{Unlock}(x)$
 ~~deadlock~~

③ Apply basic QPL and strict QPL to the following schedule.

T_3	T_4
$r_1(x)$	$r_2(x)$
$w_1(x)$	$w_2(x)$
$r_1(4)$	
$w_1(4)$	

Basic QPL

T_3	T_4
$\text{xlock}(x)$	
$r_1(x)$	
$w_1(x)$	
$\text{lock}(4)$	
$\text{unlock}(4)$	
	$\text{xlock}(x)$
	$r_2(x)$
	$w_2(x)$
$r_1(4)$	
$w_1(4)$	
$\text{unlock}(4)$	

Strict QPL

T_3	T_4
$\text{xlock}(x)$	
$r_1(x)$	
$w_1(x)$	
$\text{lock}(4)$	
	$\text{xlock}(x)$
	w_{wait}
	$r_3(4)$
	$w_3(4)$
	commit
	$\text{xlock}(x)$
	$r_2(x)$
	$w_2(x)$
	commit

4 Consider a schedule of 4 transactions T_1, T_2, T_3, T_4 and its corresponding log entry at the point of a system crash.

[start-transaction, T_1]

Initial values : $A = 30, B = 10, C = 30, D = 20$

[read-item, T_1, A]

[read-item, T_1, D]

[write-item, $T_1, D, 20, 25$]

[commit, T_1]

checkpoint

[start-transaction, T_2]

[read-item, T_2, B]

[start-transaction, T_4]

[read-item, T_4, D]

[write-item, $T_4, D, 25, 15$]

[write-item, $T_2, B, 12, 18$]

[commit, T_2]

[start-transaction, T_3]

[write-item, $T_3, C, 30, 40$]

[read-item, T_4, A]

[write-item, $T_4, A, 30, 20$]

[commit, T_4]

[read-item, T_3, D]

[write-item, $T_3, D, 15, 25$] \rightarrow System Crash

a. Use the deferred update protocol & apply the necessary redo operations

Solution : REDO type entries contain only AFIM values in the write-item operations:

[start-transaction, T₁]

[write-item, T₁, D, 05]

[commit, T₁]

[checkpoint]

[start-transaction, T₂]

[start-transaction, T₄]

[write-item, T₄, D, 15]

[write-item, T₂, B, 18]

[commit, T₂]

[start-transaction, T₃]

[write-item, T₃, C, 40]

[write-item, T₄, A, 20]

[commit, T₄]

[write-item, T₃, D, 05]

For Redo operations, for everything after the checkpoint,

Maintain 2 lists

REDO LIST

T₄, T₂

ACTIVE LIST

T₃ → push again

↳ is uncommitted

rewrite items in the
same order.

[write-item, T₄, D, 15]

[write-item, T₂, B, 18]

[write-item, T₄, A, 20]

B. Remove the [commit, T₂] entry. Use the immediate
update policy to make the necessary undo/redo operations

commit list

T₄

redo T₄ operations
in the same order

[write-item, T₄, D, 15²⁵]

[write-item, T₄, A, 30, 20]

active list

T₂, T₃

undo operations in the backward

order

[write-item, T₃, D, 15, 25]

[write-item, T₃, C, 30, 40]

[write-item, T₂, B, 12, 18]

Transaction Processing

* Transaction: provides a mechanism for describing the logical units of a database processing

Classification of database systems

(i) single-user system: at most one user at a time

(ii) multi-user system: many users can access the system concurrently

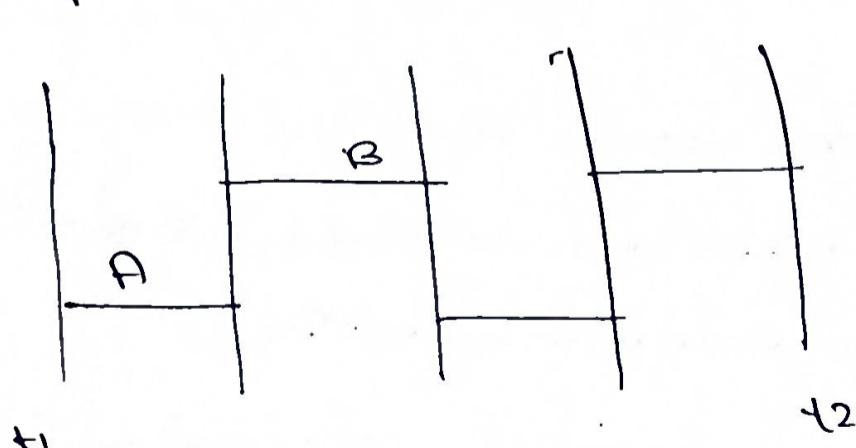
Interleaved and Parallel Processing

A. Interleaved Processing

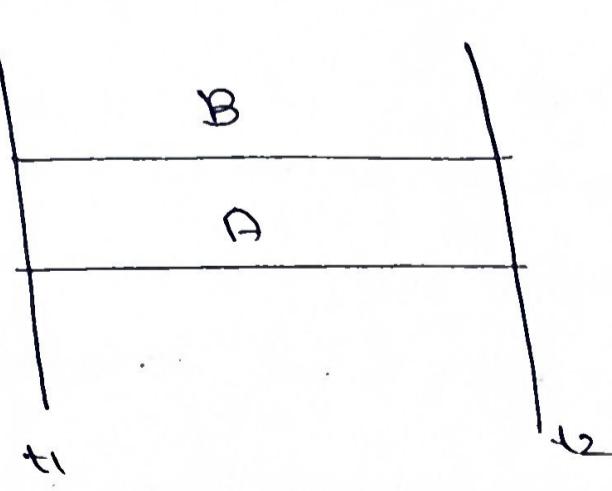
- concurrent execution of processes is interleaved in a single CPU
- keeps CPU busy when reading I/O operations, i.e. minimum idle time during I/O
- prevents long processes from delaying other processes

B. Parallel Processing

- processes are concurrently executed in multiple CPUs



Interleaved Processing



Parallel Processing

* Granularity

- granularity is the size of the data item
- granularity of data = a field, record or a whole block
- Transaction processing concepts are independent of granularity and apply to data items in general

* Read and Write Operations on Databases

- ① Read-item (X): reads a database item named X into a program variable X

Steps

- find address of disk block that contains memory item X
- copy that disk block into a buffer in main memory
- copy item X from buffer to program variable X

- ② Write-item (X): writes the value of program variable X into the database item named X.

Steps

- find address of disk block that contains item X
- copy that disk block into a buffer in main memory (if that disk block is not already in some main memory buffer)
- copy item X from the program variable named X into the correct location on the buffer
- store the updated block from buffer back to disk

* Transaction Notation

(3)

Ex1 T₁

read-item(X);
X' = X - n
write-item(A);
read-item(Y);
Y = Y + N;
write-item(Y);

Ex2 T₂

read-item(X);
X' = X + M;
write(X);
item

T₁: ~~r₁(x)~~

b₁; r₁(X); w₁(X); r₁(Y);
w₁(Y); e₁

b₂;
T₂: r₂(X); w₂(X); e₂;

* Concurrency control & database recovery mechanism → deal with the

database commands in a transaction.

→ uncontrolled concurrent execution would lead to database inconsistency

* Sources of Database Inconsistency

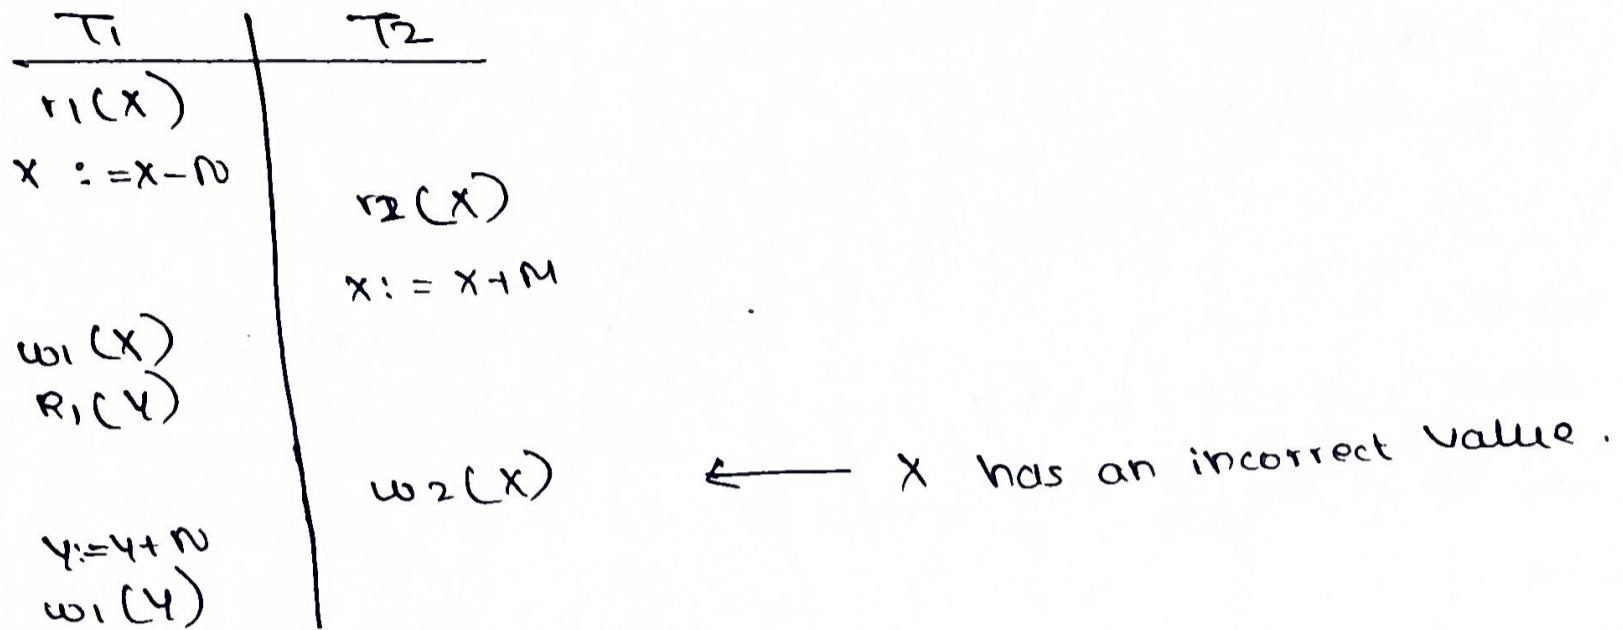
→ uncontrolled execution of database transactions in a multi-user environment can lead to db inconsistency.

Source of db inconsistency

- (i) Lost Update Problem
- (ii) Dirty Read Problem
- (iii) Incorrect summary problem
- (iv) Unrepeatable Read problem

① Last Update Problem

→ This occurs when two transactions that access the same database items have their operations interleaved in a way that makes the value of the same database item incorrect.



Suppose $X = 100$

$N = 5$

and $M = 4$

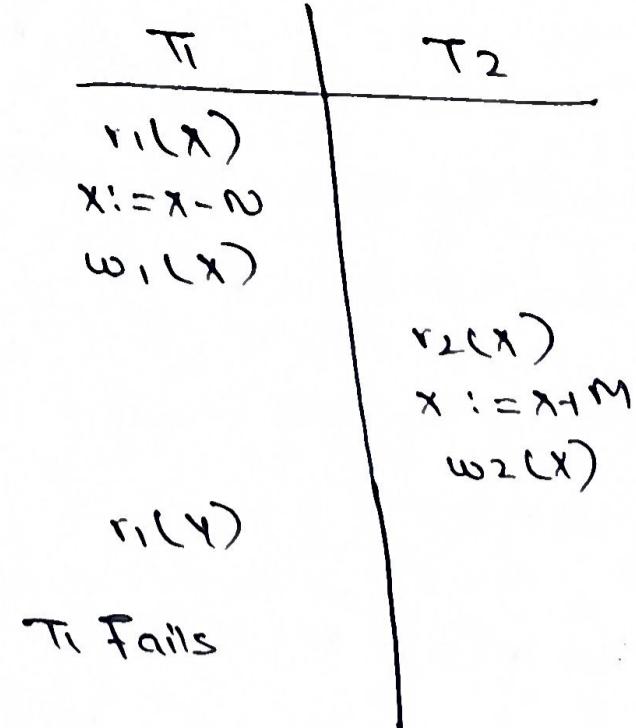
The final value of X should be 99
However, T_2 reads the value of X before T_1 changes it in the database \Rightarrow The value of X would be 104 instead, because the update in T_1 that removed 5 was lost.

② The Temporary Update / Dirty Read Problem

→ This occurs when one transaction updates a database item and then the transaction fails for some reason.

→ The updated item is accessed by another transaction before it is changed back to its original value.

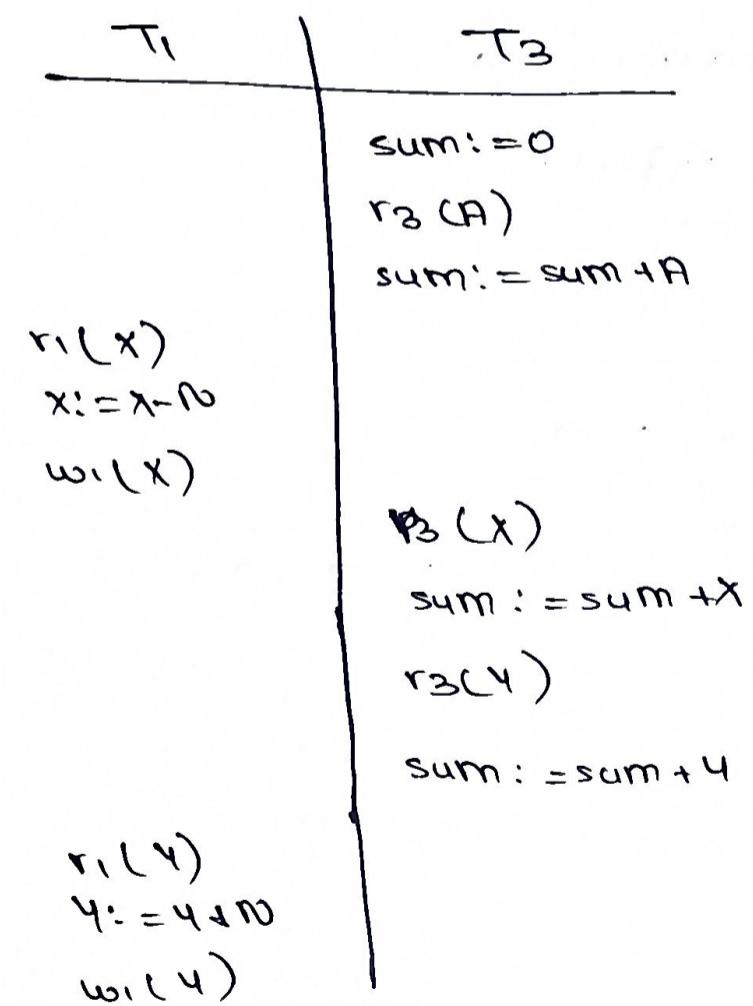
(5)



Transaction T_1 fails and must change the value of X back to its old value, meanwhile T_2 has read the temporary incorrect value of X .

③ Incorrect Summary Problem

→ If one transaction is calculating an aggregate summary function on a number of records, while other transactions are updating some of these records, the aggregate fn. may calculate some values before they are updated and others after they are updated.

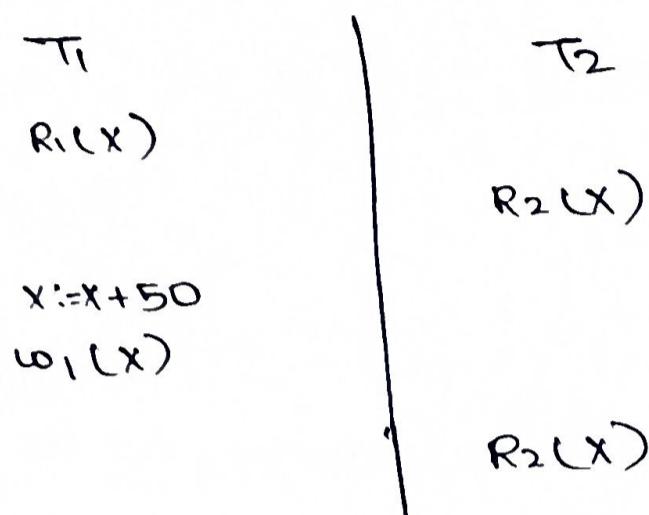


T_3 reads after N is subtracted
reads
and before N is added to Y .

⇒ The sum is off by N .

④ Unrepeatable Read Problem

- When a transaction reads the same item twice and the item is changed by another transaction T between the 2 reads.
- T receives a different for 2 reads of the same item



if $X = 100$

→ The initial R₂(X) reads 100,
but the second R₂(X) reads
150.

* Why is recovery needed?

- If the transaction fails after executing some of the operations but before executing all of them, the operations already executed must be undone and must have no lasting effect.

* Types of Failures

classified as transaction, system and media failures

(i) Computer failure (system crash): hardware | software | network error during transaction

(ii) Transaction or system error: operations that cause overflow, division by 0, logical programming errors etc.

(iii) Local errors or exception conditions detected by the transaction: cancellation of transaction due to lack of data (insufficient fund)

(iv) concurrency control enforcement : transactions aborted

due to deadlock, timeout, serializability violations

(v) disk failure : read / write malfunction

(vi) physical problems / catastrophe : power failure, theft etc.

DBMS has a Recovery Subsystem to protect db against system failures

* Transaction States

→ A transaction is an atomic unit of work that is either completed in its entirety or not done at all.

→ The different transaction states are:

(i) Active State : Transaction goes into active state immediately after it starts execution

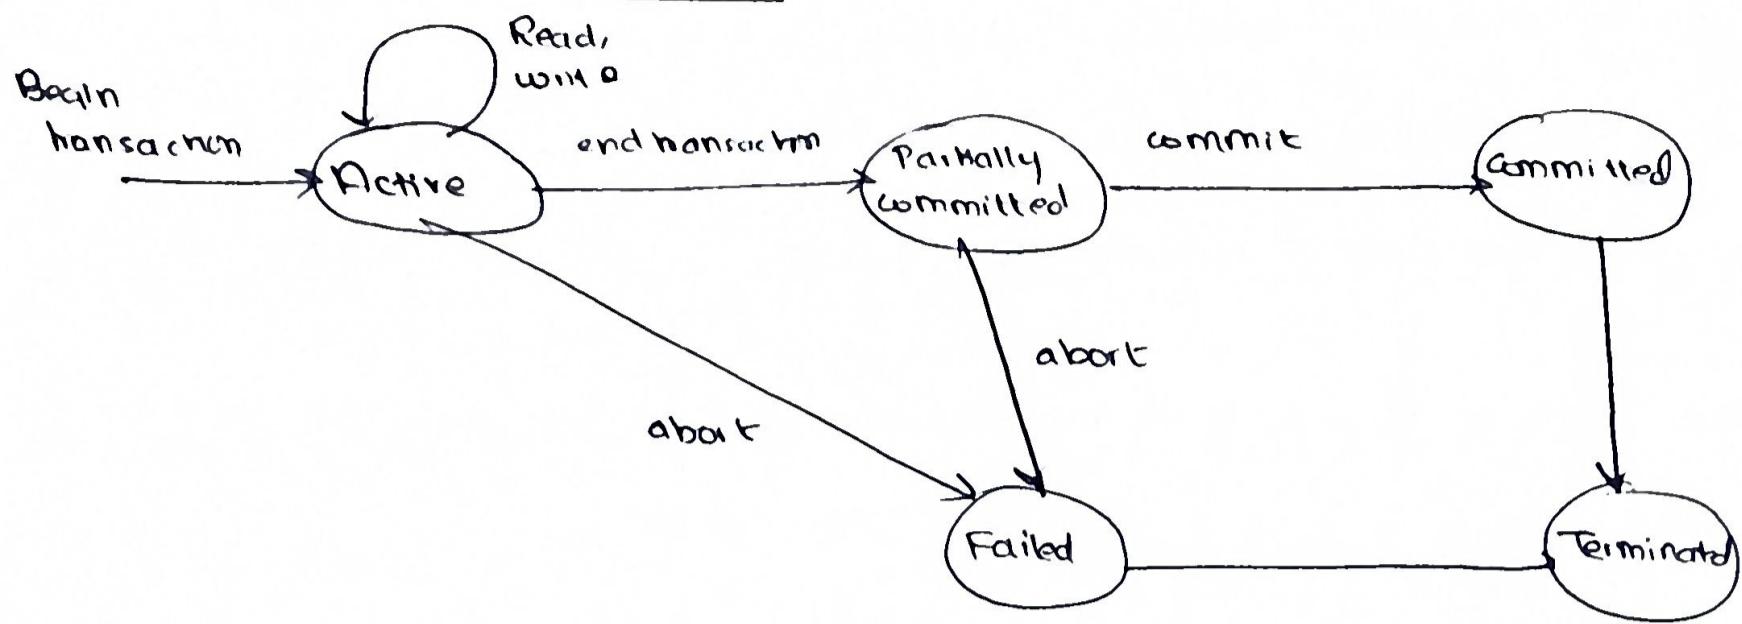
(ii) Partially Committed State : when the transaction ends, it moves to the partially committed state

(iii) Committed State : If the transaction reaches its commit point, the transaction enters the committed state

(iv) Failed State : If the transaction is aborted during its active state, it enters into the failed state.

(v) Terminated State corresponds to the transaction leaving the system

* State Transition Diagram



The recovery manager keeps track of the following operations

- (i) begin-transaction: marks beginning of transaction execution
- (ii) Read/write : the specified read/write operations on the db items that are executed as part of a transaction
- (iii) end-transaction: marks that read & write ops have completed.
- (iv) commit-transaction : signals a successful end of transaction, any changes can be safely committed to the db.
- (v) rollback/abort: transaction ended unsuccessfully, any changes & effects must be undone.

* System Log

- To be able to recover from transaction failures, the system maintains a log.
- The log keeps track of all transaction operators that affect the values of db items, as well as info needed to permit recovery

- The log is a sequential, append-only file
- The log is kept on a disk, so it is not affected by any type of failure, also periodically backed up to archival storage.

* Contents of Log File

A log file has the following contents:

T refers to a unique transaction ID

- (i) [start-transaction, T] : records that transaction T has started execution
- (ii) [write-item, T, X, old-value, new-value] : records that T has changed the value of database item X from old-value to new-value
- (iii) [read-item, T, X] : records that transaction T has read the value of database item X
- (iv) [commit, T] : records that transaction T has completed successfully, and affirms that its effect can be committed to the db
- (v) [abort, T] : records that transaction T has been aborted.

* UNDO & REDO Operations

UNDO: reset all items changed by a write operation of T to their old-values by tracing backwards through the log records

REDO: by setting all items changed by a write operation of T to their new-values. Redo may be necessary when a transaction

has been written into the log, but a failure occurs before the new-values have been written to the actual database.

* Writing of Log Files

- The log file must be kept on disk.
- It is common to keep one or more blocks of the log file in the main memory buffers, until they are filled with log entries, and then to write them back to the disk only once, rather than writing to the disk every time a log entry is added.

- At the ^{time of} system crash, only the log entries that have been written back to the disk are considered in the recovery process.

Force Writing a Log : Before a transaction reaches its commit point, any portion of the log that has not been written to the disk yet must now be written to the disk. This is called force-writing the log file before committing a transaction.

* ACID Properties

- A transaction should possess several properties called ACID properties that should be enforced by the concurrency control recovery methods.

(1) Atomicity

- a transaction is an atomic unit of processing
- it is either performed in its entirety, or not performed at all
- responsibility of the transaction recovery subsystem.

- (II) Consistency/Preservation

- every transaction should execute w/o the interference of other transactions
- transaction must take db from one consistent state to another
- programmer's responsibility or the DBMS module to enforce integrity constraints

(III) Isolation

- transaction should not make its update visible to other transactions until it is committed
- solves the temporary update problem & eliminates cascading rollbacks

Level 0 : no dirty reads

Level 1 : no lost updates

Level 2 : no lost updates & no dirty reads

Level 3 : Level 2 + repeatable reads

- enforced by concurrency control subsystem of DBMS

(IV) Durability or permanency

- once a transaction changes the db and the changes are committed these changes must never be lost because of subsequent failure
- responsibility of the recovery subsystem

Recovery protocols enforce atomicity & durability

Example Transfer \$50 from account a to b

read(a) : $a := a - 50$; write(a)

read(b) : $b := b + 50$; write(b)

Atomicity: a must not be debited w/o crediting b

Consistency: sum of a & b remains constant

Isolation: if another transaction reads a and b, it must see one of the 2 states: before or after the transaction, otherwise $a+b$ is wrong

Durability: if b is notified of credit, it must persist even if the database crashes.

Schedules

* Schedule = history : ordering of operations of transactions

* Conflicting Operations in a Schedule:

(i) belong to different transactions

: can be

read-write

or write-write

(ii) access same data item X

(iii) at least one operation is a write-item (X)

Ex. Find conflicting operations

S: $r_1(x), r_2(x), w_1(x), r_1(y), w_2(y), w_1(y)$;

Ans $r_1(x), w_2(x)$
 $r_2(x), w_1(x)$
 $w_1(x), w_2(x)$

complete schedule

operations in S are those in $T_1, T_2, T_3 \dots$

relative ordering is the same

for conflicting operations - one must happen before the other.

conflicting operations can happen because of a change in order resulting in a different outcome

(i) read-write conflict $r_1(x); w_2(x) \neq w_2(x); r_1(x)$

(ii) write-write conflict $w_1(y); w_2(y) \neq w_2(y); w_1(y)$

* Characterizing Schedules based on Recoverability

Recoverable Schedule : make sure that schedule w/ write ops (from which another schedule has read from) commits first

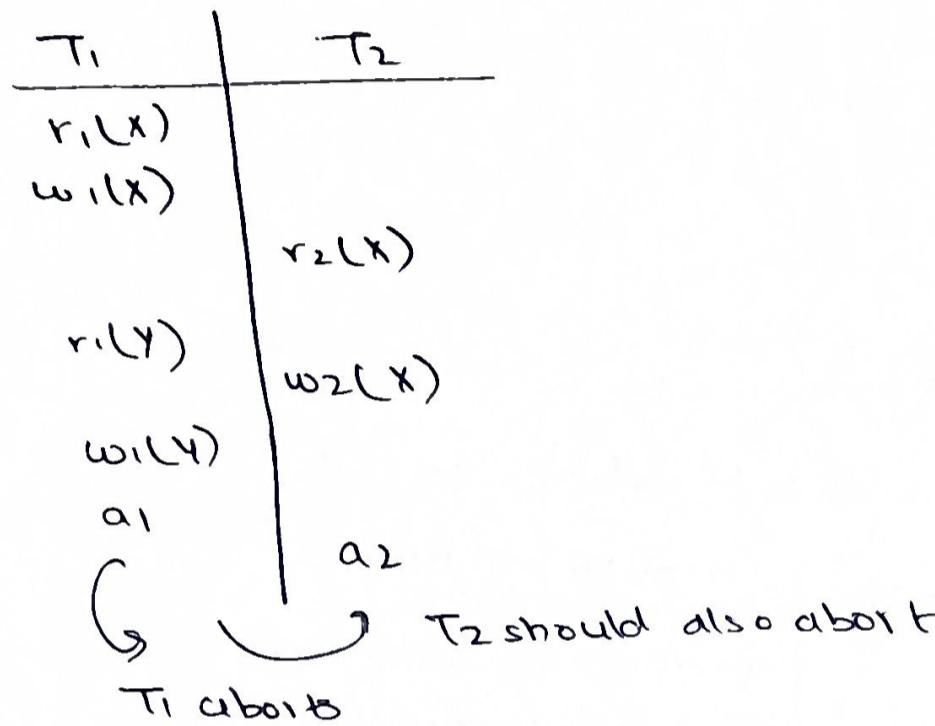
aka w commits first

r commits next

* Cascading Rollback

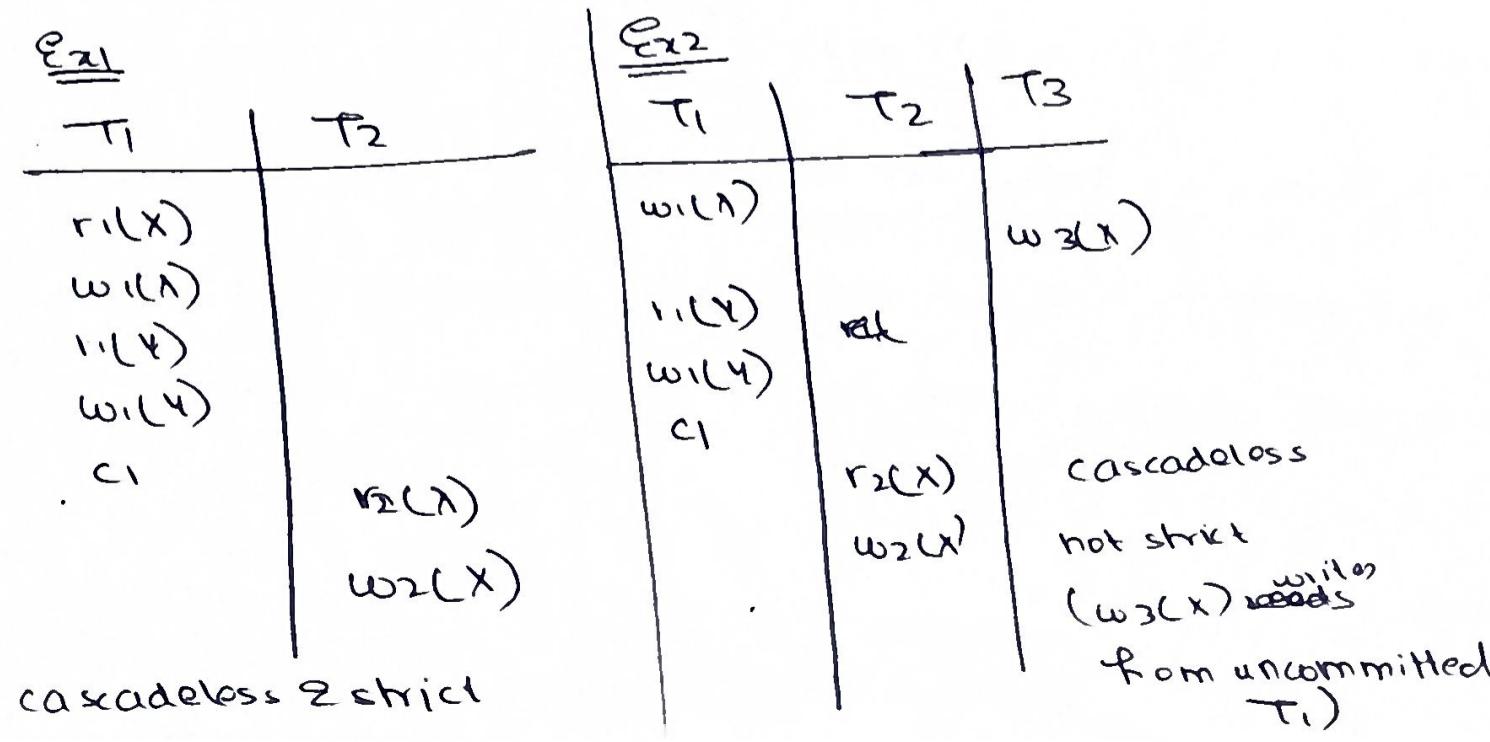
→ an uncommitted transaction has to be rolled back because it read an item from a transaction that failed

S: $i_1(x); w_1(x); r_2(x); r_1(y); w_2(x); w_1(y); a_1; a_2$



* Cascadeless Schedule: read items only from committed transactions (blind writes allowed)

* Strict Schedule: read and write only allowed from committed transactions



Ex 3

T ₁	T ₂	T ₃
r ₁ (x)		
w ₁ (x)		neither cascadeless nor strict
r ₁ (y)		T ₃ should commit,
w ₁ (y)		then only T ₂ can read (cascadeless)
q		
	w ₂ (x)	
	r ₂ (x)	for T ₂ to write, T ₃ should
	w ₂ (x)	commit (strict)

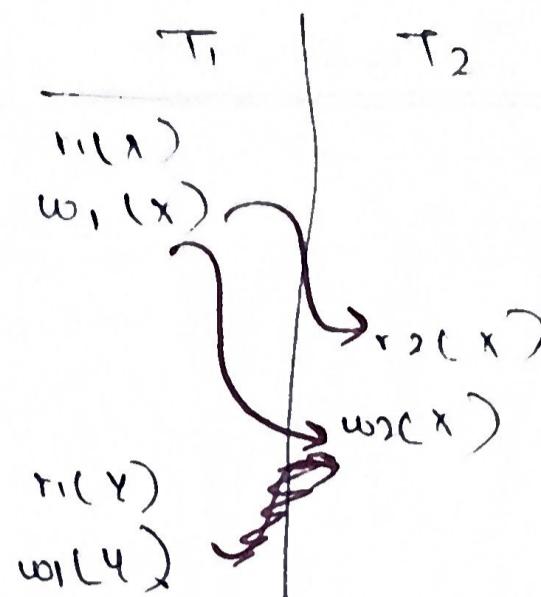
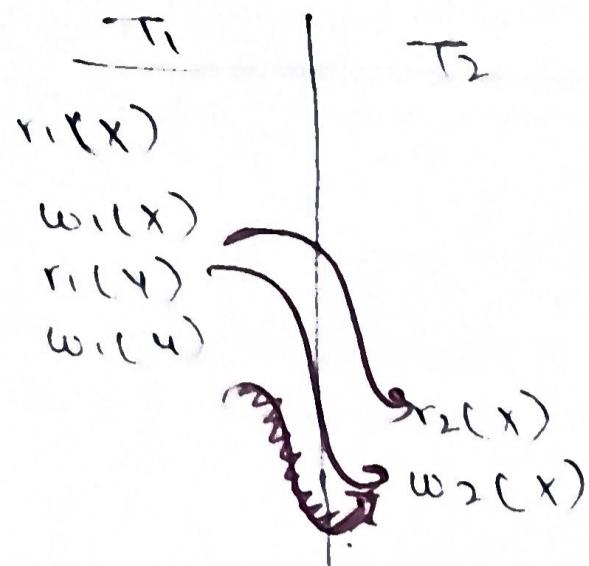
* Serializability Theory → determines which schedules are correct

A schedule S is said to be serializable if it is equivalent to some serial schedule of the same n transactions.

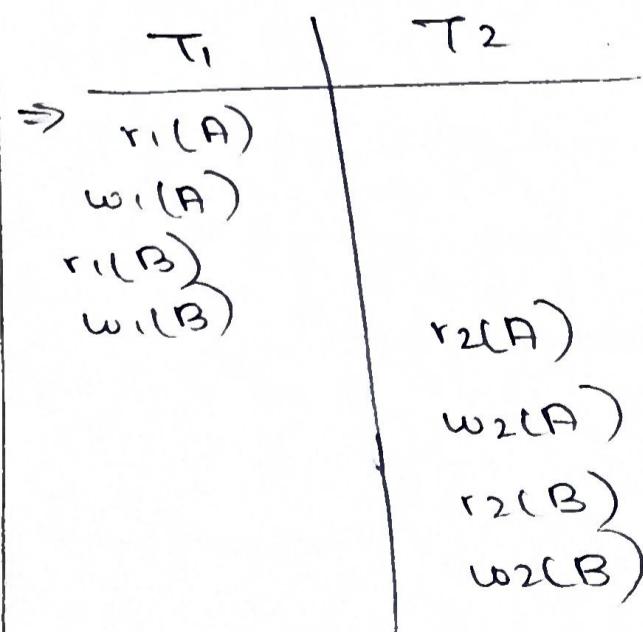
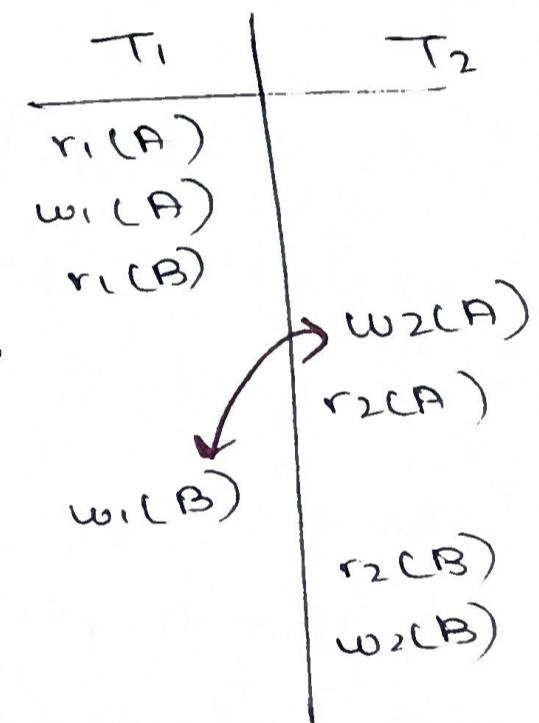
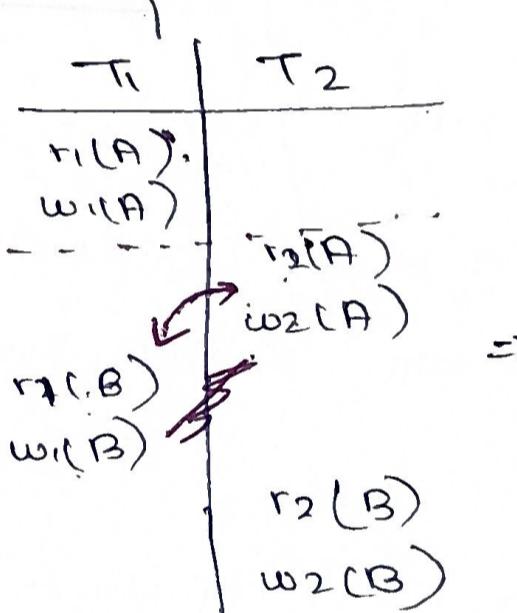
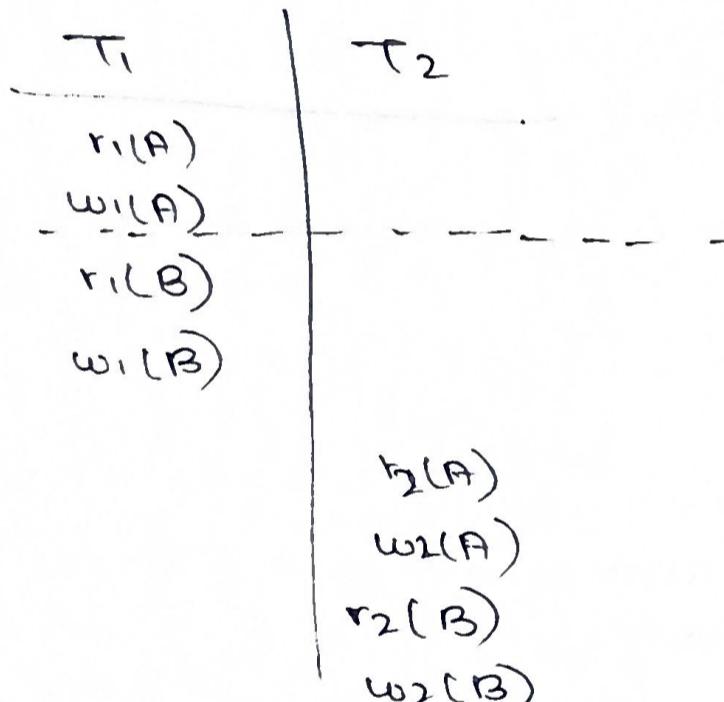
considered equivalent when

- (i) result equivalent → produces same result at the final state.
- (ii) conflict equivalent → ordering of conflict equivalence is the same done using
 - (i) ordering conflict ops
 - (ii) swapping non-conflict ops
 - (iii) precedence graph
- (iii) view equivalent

Ordering of Conflict Operations



Swapping of non-conflict operations



View Serializable

UNIT 4

PRACTICE QUESTIONS ON SERIALIZABILITY

① Classify the following schedules as cascadeless, strict etc.

A. Schedule C: $r_1(x); w_1(x); r_1(y); w_1(y); c_1; r_2(x); w_2(x)$

T_1	T_2
$r_1(x)$	
$w_1(x)$	
$r_1(y)$	
$w_1(y)$	
c_1	
	$r_2(x)$
	$w_2(x)$

cascadeless \Rightarrow commit ~~after write~~^{before read} operation

strict \rightarrow commit before read & write

B. Schedule D: $r_1(x); w_1(x); w_3(x); r_1(y); w_1(y); c_1; r_2(x); w_2(x)$

T_1	T_2	T_3
$r_1(x)$		
$w_1(x)$		
$r_1(y)$		
$w_1(y)$		
c_1		
	$r_2(x)$	
	$w_2(x)$	
		$w_3(x)$

cascadeless \Rightarrow read happens after a commit

not strict $\Rightarrow T_3$ not committed, $w_3(x)$ is written before that (blind writes not allowed)

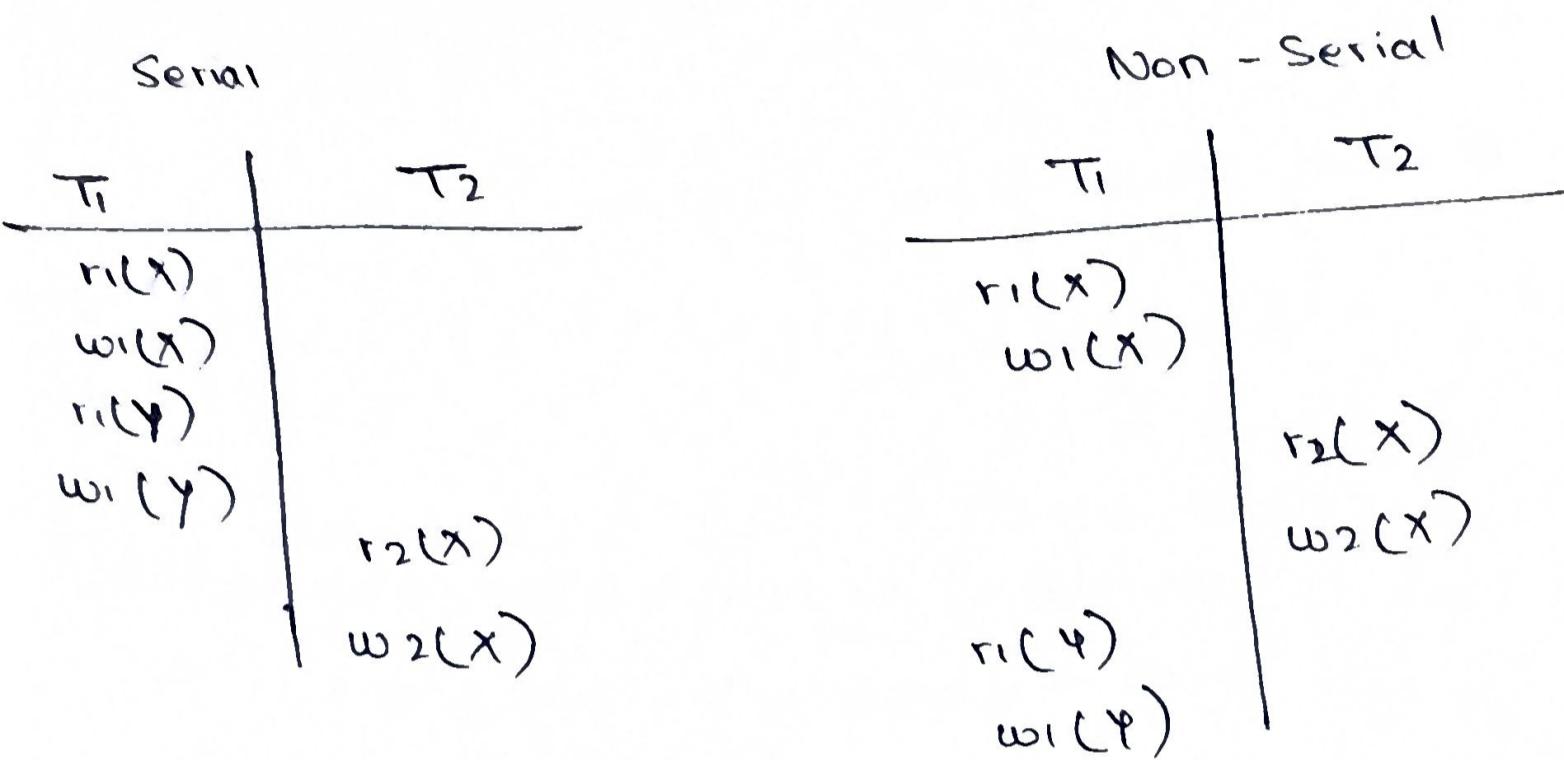
C. Schedule E: $r_1(x); w_1(x); r_1(y); c_1; w_3(x); r_2(x); w_2(x)$

T_1	T_2	T_3
$r_1(x)$		
$w_1(x)$		
$r_1(y)$		
c_1		
	$r_2(x)$	
	$w_2(x)$	
		$w_3(x)$

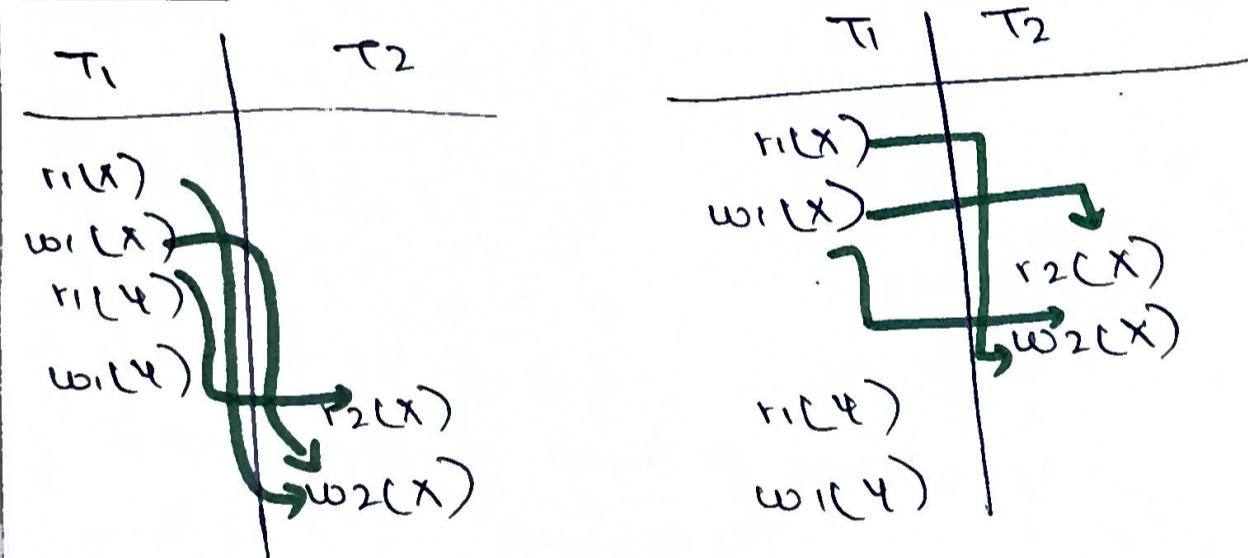
not cascadeless $\Rightarrow T_3$ not committed before write operation

not strict $\Rightarrow T_3$ not committed before $w_2(x)$

② Check if the non-serial schedule is serializable based on the ordering of conflict operations



Solution



Ordering of conflict operations

- $r_1(x) \rightarrow w_2(x)$
- $w_1(x) \rightarrow r_2(x)$
- $w_1(x) \rightarrow w_2(x)$

are the same in the serializable & non-serializable schedules.

\therefore The schedules are conflict operation equivalent.

③ Check if the 2 schedules are conflict equivalent by ③⁽³⁾
the swapping of non-conflict operations

Serial		Non-serial	
T ₁	T ₂	T ₁	T ₂
r ₁ (A)			r ₁ (A)
w ₁ (A)			w ₁ (A)
r ₁ (B)			r ₂ (A)
w ₁ (B)			w ₂ (A)
	r ₂ (A)	r ₁ (B)	
	w ₂ (A)	w ₁ (B)	
	r ₂ (B)		r ₂ (B)
	w ₂ (B)		w ₂ (B)

Solution

T ₁	T ₂
r ₁ (A)	
w ₁ (A)	
r ₁ (B)	r ₂ (A)
w ₁ (B)	w ₂ (A)
	⇒
r ₁ (B)	r ₁ (A)
w ₁ (B)	w ₂ (A)
r ₂ (B)	r ₂ (A)
w ₂ (B)	w ₂ (B)

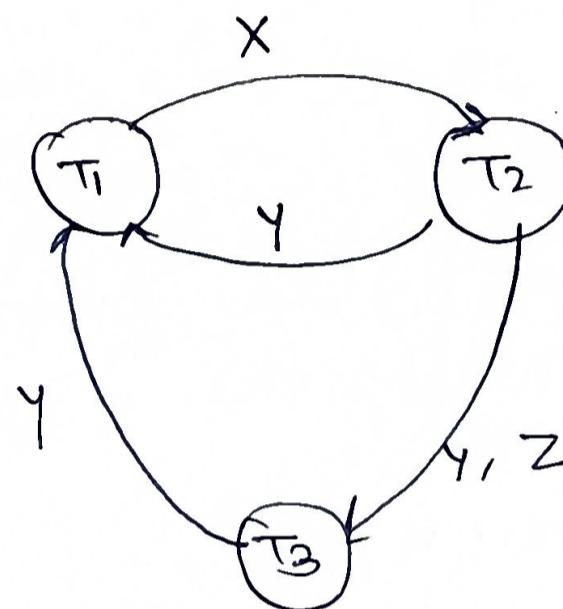
T ₁	T ₂
r ₁ (A)	
⇒ w ₁ (A)	
r ₁ (B)	
w ₁ (B)	
	r ₂ (A)
	w ₂ (A)
	r ₂ (B)
	w ₂ (B)

, which is the equivalent serial schedule.

∴ It is conflict-serializable.

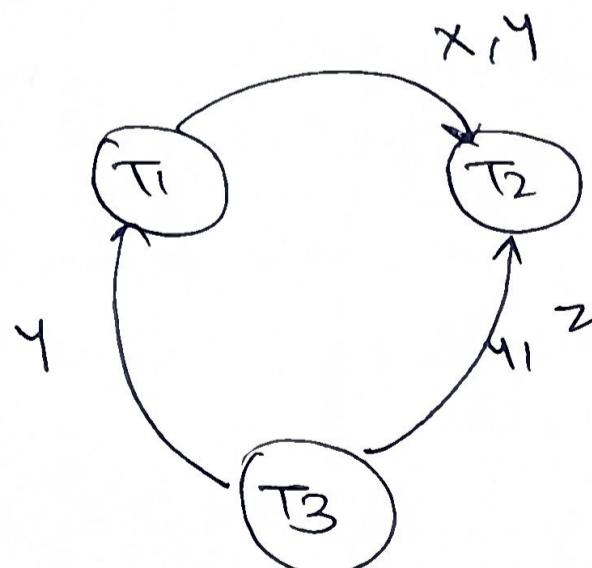
④ Draw the precedence graphs for the following schedules, and check whether they are conflict serializable or not.

	T_1	T_2	T_3
		$r_2(2)$	
		$r_2(4)$	
		$w_2(4)$	
			$r_3(4)$
			$r_3(2)$
			$w_3(4)$
			$w_3(2)$
		$r_2(x)$	
			$w_3(x)$
		$r_1(4)$	
		$w_1(4)$	
		$w_2(x)$	



not serializable - there exists a cycle

	T_1	T_2	T_3
			$r_3(4)$
			$r_3(2)$
			$w_3(4)$
			$w_3(2)$
		$r_2(2)$	
			$w_3(x)$
		$r_1(4)$	
		$w_1(4)$	
		$r_2(4)$	
		$w_2(4)$	
		$r_2(x)$	
		$w_2(x)$	



no cycles \Rightarrow serializable
Serializable schedule

$\rightarrow T_3 \rightarrow T_1 \rightarrow T_2$

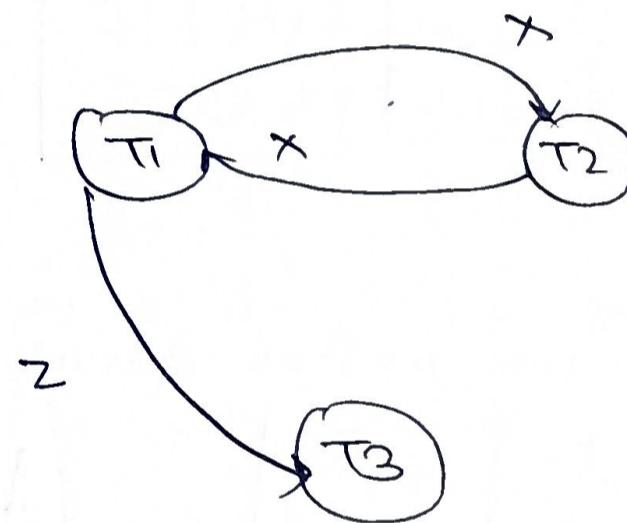
(3) Consider the following schedule. Check whether it is serializable or not.

$r_1(x); r_3(y); r_2(x); w_1(x); w_2(x); w_3(y);$
 $r_1(z); w_1(z); r_3(z); w_3(z)$

If not serializable, swap $r_2(x)$ and $w_1(x)$. Swap non-conflicting operations to generate the equivalent serial schedule

Solution

T_1	T_2	T_3
$r_1(x)$		$r_3(y)$
	$r_2(x)$	
$w_1(x)$		$w_2(x)$
		$w_3(y)$
$r_1(z)$		
$w_1(z)$		
	$r_3(z)$	
		$w_3(z)$



There exists a cycle:

\Rightarrow not serializable

swap $r_2(x)$ and $w_1(x)$

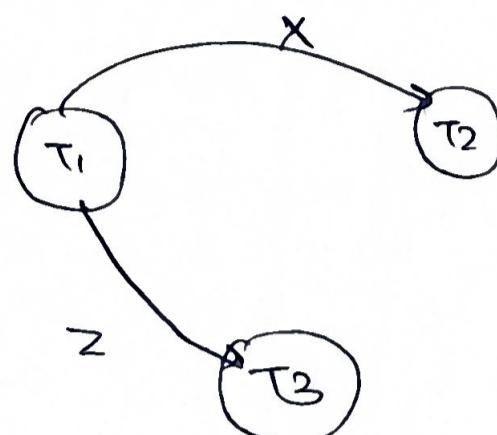
The schedule is now

$r_1(x); r_3(y); w_1(x); r_2(x); w_2(x); w_3(y);$
 $r_1(z); w_1(z); r_3(z); w_3(z)$

$r_1(x)$; $r_3(4)$; $w_1(x)$; $r_2(x)$; $w_2(x)$; $w_3(4)$; $r_1(2)$;
 $w_1(2)$; $r_3(2)$; $w_3(2)$

T_1	T_2	T_3
$r_1(x)$		
$w_1(x)$		
	$r_2(x)$	
	$w_2(x)$	
		$w_3(4)$
$r_1(2)$		
$w_1(2)$		
	$r_3(2)$	
		$w_3(2)$

Precedence Graph



no-cycle \Rightarrow serializable

order: $T_1 \rightarrow T_2 \rightarrow T_3$

or $T_1 \rightarrow T_3 \rightarrow T_2$

Swapping non-conflict operations to obtain a serial schedule

T_1	T_2	T_3
$r_1(x)$		
$w_1(x)$		
	$r_2(x)$	
	$w_2(x)$	
		$w_3(4)$
$r_1(2)$		
$w_1(2)$		
	$r_3(2)$	
		$w_3(2)$

T_1	T_2	T_3
$r_1(x)$		
$r_1(2)$	$r_2(x)$	$r_3(4)$
$w_1(x)$		
	$r_2(x)$	$r_3(4)$
	$w_2(x)$	
		$w_3(4)$
$r_2(x)$		
$w_1(2)$		
		$r_3(4)$
		$r_3(2)$
		$w_3(2)$

T ₁	T ₂	T ₃
r ₁ (x)		
w ₁ (2)		
w ₁ (x)		
w ₁ (2)		
	w ₂ (x)	
		w ₃ (4)
		r ₃ (4)
	r ₂ (x)	
		r ₃ (2)
		w ₃ (2)

(7)

T ₁	T ₂	T ₃
r ₁ (x)		
w ₁ (2)		
w ₁ (x)		
w ₁ (2)		
	w ₂ (x)	
		r ₂ (x)
		r ₃ (4)
		w ₃ (4)
		r ₃ (2)
		w ₃ (2)

swap w₂(x) and r₃(2)

and r₂(x) and w₃(2)

to get T₁ → T₃ → T₂

T₁ → T₂ → T₃

equivalent serial
schedule.

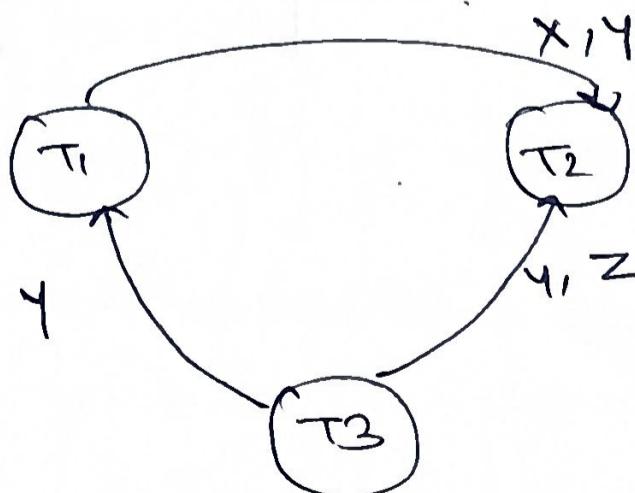
- (6) Consider the following schedule, check if it is serializable or not. If it is, generate the equivalent serial schedule.

r₃(4); r₃(2); r₁(x); w₁(x); w₃(4); w₃(2); r₂(2);
r₁(4); w₁(4); r₂(4); w₂(4); r₂(x); w₂(x);

Solution

T ₁	T ₂	T ₃
r ₁ (x)		r ₃ (4)
w ₁ (x)		r ₃ (2)
	w ₃ (4)	
r ₂ (2)		w ₃ (2)
r ₁ (4)		
w ₁ (4)		
r ₂ (4)		
w ₂ (4)		
r ₂ (x)		
w ₂ (x)		

Precedence Graph



no cycles \Rightarrow serializable

order of serial schedule: T₃ → T₁ → T₂

Swapping of non-conflicting operations to obtain a serial schedule

T_1	T_2	T_3
$r_1(x)$		$r_3(4)$
$w_1(x)$		$r_3(2)$
		$w_3(4)$
		$w_3(2)$
$r_1(x)$	$r_2(4)$	
$w_1(x)$	$w_2(4)$	
$r_2(x)$		
$w_2(x)$		

\Rightarrow

T_1	T_2	T_3
$r_1(x)$		$r_3(4)$
$w_1(x)$		$r_3(2)$
	$r_2(4)$	
	$w_2(4)$	
	$r_2(x)$	
	$w_2(x)$	

T_1	T_2	T_3
$r_1(x)$		$r_3(4)$
$w_1(x)$		$r_3(2)$
$w_1(x)$		$w_3(4)$
$r_1(x)$		$w_3(2)$
	$r_2(4)$	
$w_2(4)$		
$w_2(4)$		
$r_2(x)$		
$w_2(x)$		

\Rightarrow

T_1	T_2	T_3
$r_1(x)$		$r_3(4)$
$w_1(x)$		$r_3(2)$
$w_1(x)$		$w_3(4)$
$r_1(x)$		$w_3(2)$
	$r_2(4)$	
	$w_2(4)$	
	$r_2(x)$	
	$w_2(x)$	

Steps to identify View Serializability

(9)

To check if a schedule is view serializable ,

check for (i) initial read

(ii) final write

order of update operations must also be the same .

Alternatively,

first check if the schedule is conflict - serializable

If conflict - serializable \Rightarrow view serializable

If not conflict - serializable \Rightarrow may or may not be view
serializable

(i) check for blind writes : no blind writes

\uparrow

a write operation without
a read operation

\Rightarrow not

view serializable

If there are blind writes \Rightarrow may or

may not be view serializable

~~\Rightarrow draw dependency graph~~

write down initial read, update
& final read values

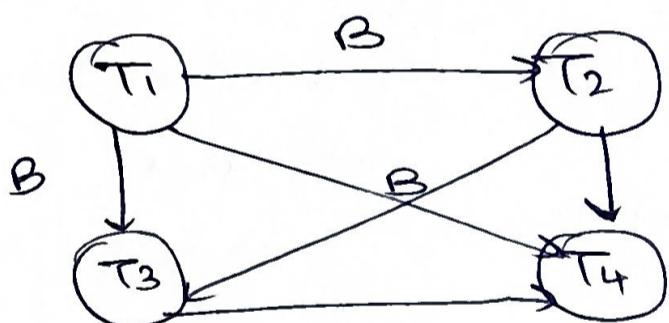
Problems

- ① Check if the following schedule is view serializable or not

T ₁	T ₂	T ₃	T ₄	
R(A)				(OR)
	R(A)			
		R(A)		
			R(A)	
W(B)				
	W(B)			
		W(B)		
			W(B)	

A	B	
T ₁ , T ₂ T ₃ , T ₄	-	
-	T ₁ , T ₂ T ₃ , T ₄	
-	-	T ₁ , T ₂ T ₃ , T ₄

Solution: (i) check for conflict serializability



no cycles \Rightarrow conflict
serializable

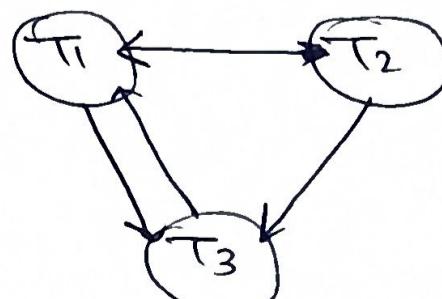
\Rightarrow view serializable

order: T₁ \rightarrow T₂ \rightarrow T₃ \rightarrow T₄

- ② Check whether the given schedule is view serializable or not

T ₁	T ₂	T ₃
R(A)		
	R(A)	
		W(A)
W(A)		

conflict serializability



not conflict
serializable

\Rightarrow may or may not be
view serializable

check for blind writes

4

$w_3(A)$ is a blind write

\Rightarrow may or may not be view serializable

Finding Dependencies

	A
Initial Read	T_1, T_2
Update	T_1, T_3
Final write	T_1, T_3

$$\begin{array}{l} T_1 \rightarrow T_3 \\ T_2 \rightarrow T_1 \\ T_3 \rightarrow T_1 \end{array} \quad \left. \begin{array}{l} \\ \\ \end{array} \right\} \text{conflicting}$$

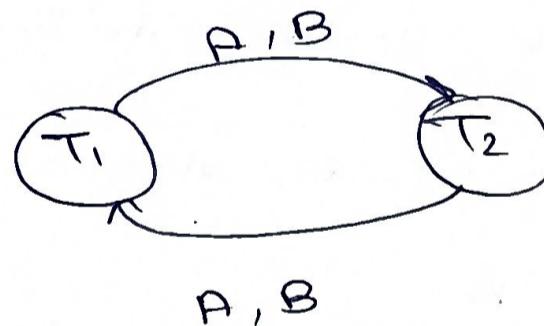
not view-serializable

③ Check whether the given schedule is view serializable or not.

	T_1	T_2
	R(A)	
	R(A)	
w(A)		w(A)
		w(A)
R(B)		R(B)
		R(B)
w(B)		
		w(B)

Solution

Check for conflict serializability



not conflict serializable \Rightarrow may or may

not be view serializable

(ii) check for blind writes - no blind writes \Rightarrow not view serializable

	A	B
Initial Read	T_1, T_2	T_1, T_2
Update	T_1, T_2	T_1, T_2
Final Write	T_2	T_2

$$\begin{array}{l} T_1 \rightarrow T_2 \\ T_2 \rightarrow T_1 \end{array} \quad \left. \begin{array}{l} \\ \end{array} \right\} \text{conflicting}$$

4

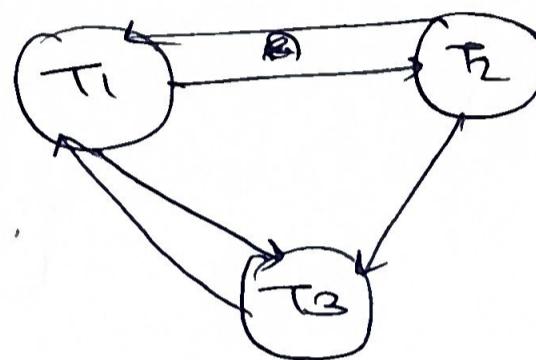
Check if the given schedule is view serializable or not. If:

Yes, give the serial schedule.

S: $R_1(A), W_2(A), R_3(A), W_1(A), W_3(A)$

T_1	T_2	T_3
$R_1(A)$		
	$W_2(A)$	
		$R_3(A)$
$W_1(A)$		
		$W_3(A)$

conflict serializability



not conflict serializable \Rightarrow may or may not be view serializable

Blind write check: $w_2(A)$ in T_2 is a blind write

\Rightarrow may or may not be view serializable.

	A
Initial Read	T_1, T_3
Final Update	$T_1, \cancel{T_2}, T_3$ T_1, T_2
Final Write	T_1, T_2, T_3

dependency

$T_1 \rightarrow T_2 \rightarrow T_3$

T_1 reads before T_2

$\Rightarrow T_1 \rightarrow T_2$

T_3 updates last

$\Rightarrow T_3$ executes last

$\Rightarrow T_1 \rightarrow T_2 \rightarrow T_3$

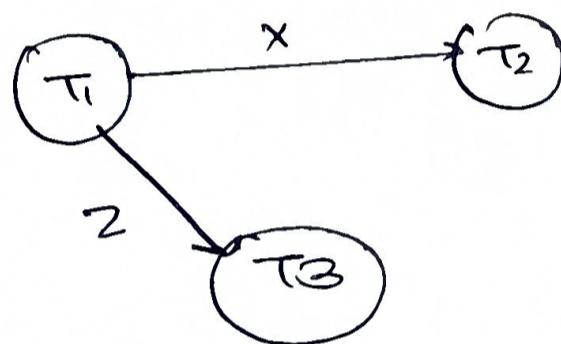
\Rightarrow |view serializable|

Check if the given schedule is view serializable

$S_1: r_1(x); r_2(4); w_1(x); r_2(x); w_2(x); w_3(4),$
 $r_1(z); w_1(z); r_3(z); w_3(z)$

	T_1	T_2	T_3
	$r_1(x)$		
		$r_3(4)$	
	$w_1(x)$	$w_1(x)$	
			$r_2(x)$
			$w_2(x)$
			$w_3(4)$
	$r_1(z)$		
	$w_1(z)$		
		$r_3(z)$	
			$w_3(z)$

(i) check if conflict serializable



conflict serializable

\Rightarrow it is view serializable

	x	4	z
Initial Read	T_1, T_2	T_3	T_1, T_3
Update	T_1, T_2	T_3	T_1, T_3
Final Write	T_2	T_3	T_3

T_1 before T_3 and T_1 before T_2

possible serial schedules - $T_1 \rightarrow T_2 \rightarrow T_3$

or $T_1 \rightarrow T_3 \rightarrow T_2$