

Distributed Database Management Systems

DDB - a collection of multiply logically related databases distributed over a computer network.

D - DBMS - a software system that manages a distributed database, while making the distribution transparent to the user.

* Conditions to be a Distributed Database

- (i) connection of database nodes over a computer network
- (ii) logical interrelation of the connected databases
- (iii) possible absence of homogeneity among connected nodes

Note that the sites may have physical proximity or may be geographically distributed

* Characteristics :-

① Transparency

A. Distribution and Network Transparency

→ Users do not have to worry about operational details and placement of data in the network

Location Transparency - freedom of issuing command from any location

Naming Transparency - allows access to any names object from any location without additional specification

B. Replication Transparency

- stores copies of data at multiple sites, minimizes access time
- better availability, performance & reliability

C. Fragmentation Transparency

- allows to fragment a relation horizontally or vertically
- fragmentation transparency makes the user unaware of the existence of fragments.

D. Design Transparency & Execution Transparency

- how the db is designed and where a transaction executes

② Increased Reliability and Availability

Reliability - the probability that a system will fail in a given period

- a distributed system is considered reliable if it keeps delivering its services even when one or several of its software / hw components fail.

Availability - probability that the system is continuously available during a time interval (3)

Fault Tolerance: D-DBNs must recognize that faults will occur & design mechanisms to detect or remove faults before they can result in an error.

③ Scalability and Partition Tolerance

Scalability - determines the extent to which a system can expand its capacity while continuing to operate without interruption.

A. Horizontal Scalability - expanding the no. of nodes in the distributed system, should distribute some data & loads from existing nodes to the new nodes

B. Vertical Scalability - expand the capacity of the individual nodes in the system, such as expanding the storage / processing power of a node.

Partition Tolerance. - Faults cause nodes to be partitioned into groups of nodes.

→ Partition Tolerance states that the system should have the capacity to continue operating while the network is partitioned.

④ Improved Performance

- fragments database to keep data closer to where it is needed most
- reduces data management time significantly
- easier expansion
- can add new nodes anytime

* Distributed Data Storage

2 approaches - replication or fragmentation

1. Replication - system maintains multiple copies of data, stored in different sites for faster retrieval & fault tolerance
2. Fragmentation - partition relation into several fragments and store in distinct sites
3. Combination of replication & fragmentation - system maintains several identical replicas of each such fragment

Data Allocation - each fragment - or each copy of a fragment must be assigned to a particular location in the distributed system.

choice of sites and degree of replication depends on the performance & availability goals of the system & the types & frequencies of transactions submitted at each site.

Adv

* Data Replication

- (i) Availability
- (ii) Parallelism
- (iii) Reduced data transfer

Disadv

- (i) Increased cost of updates
- (ii) Increased complexity of concurrency control

* Data Fragmentation

Ques. → can be horizontal or vertical

Reasons for fragmentation

- Every relation is not a suitable unit for distribution, since each application views a subset of its relations.
- Fragmentation ⇒ parallel execution of a single query by dividing it into a set of subqueries that operate on fragments
- Increases the level of concurrency (intraquery concurrency) and thereby, the throughput

A. Horizontal Fragmentation

horizontal subset of a relation satisfying selection conditions

derived horizontal fragmentation - partitioning of a primary key relation into other secondary relations related w/ foreign keys

→ specified by σ in relational algebra

→ reconstructed using \cup

B. Vertical Fragmentation

→ a subset of columns

→ no selection condition used

→ specified by π in relational algebra

→ reconstructed using natural join.

* Correctness Rule for Fragmentation

(i) Reconstruction If a relation R is decomposed into fragments $R_1, R_2 \dots R_l$, it should be possible to define a relational operator Δ such that $R = \Delta R_i$

(ii) Distinctness -

horizontal fragmentation - If a relation R is horizontally decomposed into $R_1, R_2 \dots R_n$ and a data item d_i is in R_j , it is not in any other fragment R_k .

vertical fragmentation - If a relation R is vertically decomposed, its primary key attributes are typically repeated in all the fragments

* Types of Distributed Database Systems

(i) Homogeneous - all sites of database system have identical setup
→ underlying OS may be different
(may be a mixture of Linux, Windows)

(ii) Heterogeneous

Federated - each site runs a different database system, but the access is managed through a single conceptual schema.

Multidatabase - there is no conceptual global schema, the schema is constructed dynamically as needed by the application software.

Unit-5NoSQL

* NoSQL → Not Only SQL or NO to SQL

→ Focuses on:

- (i) semistructured data storage
- (ii) high performance
- (iii) availability
- (iv) data replication
- (v) scalability

→ NoSQL systems are used for social media, posts and tweets, weblinks, spatial data etc

* When should NoSQL be used?

→ huge amount of data

→ relationship between data is not that important

→ data changes over time and is not structured

→ constraints and joins not required at the database level

→ database is growing continuously

* Characteristics of NoSQL databases

① High Availability

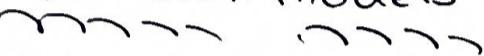
② Scalability: → determines the extent to which the system can expand its capacity without interruption

(write about horizontal & vertical scalability)

③ Replication

- data is replicated over 2 or more nodes, if one node fails, the data is still available on other nodes
- improves data availability & can improve read performance
- Eventual consistency - If no updates take place for a long time, all replicas will gradually and eventually become consistent.

Replication models



(i) Master-Slave Replication: one copy is the master or the primary copy, all write operations must be applied to the master copy and propagated to the slave copies, usually using eventual consistency.

→ Master handles writes and slaves synchronize with master.

Adv

- more read requests
- If master fails, slaves can still handle read requests
- good for datasets w/ a read-intensive dataset

Disadv

- master is a bottleneck
- failure eliminates ability to write
- inconsistency due to slow propagation of changes to slaves
- bad for datasets with heavy write traffic.

(ii) Master-Master Replication:

③

- allows reads and writes at any replicas
- all members are responsive to client data queries
- replication allows write to any node

Adv

- node failures would not affect access to data
- can easily add nodes

Disadv

- Inconsistency
- slow propagation of changes to copies on different nodes
- Two people can update different copies of the same record at the same time (write-write conflict)
- Inconsistent writes are forever.

④ Sharding of Files

- files can have millions of records
- cannot store whole file in one node
- deploy horizontal partitioning

⑤ High-Performance Data Access

- (i) Hashing
- (ii) Range partitioning - location determined via a range of key values

⑥ Does not require a schema

- semi structured and self-describing data of NoSQL systems does not require a schema.
- partial schema used to improve storage efficiency - JSON

⑦ Less Powerful Query Languages

- functions and operations as an API
- reading and writing accomplished by calling appropriate operations

CRUD operations



Create read update delete

SC RUD



search

⑧ Versioning

- provides storage of multiple versions of data items w/ timestamps

* Types of NoSQL Systems

- (i) Document-based
- (ii) NoSQL - key-value stores
- (iii) Column based
- (iv) Graph based

* CAP Theorem

→ refers to 3 desirable properties

(i) Consistency - nodes have same copies of replicated data

(ii) Availability - guarantees that every request receives a response (successful or failed)

(iii) Partition Tolerance - system continues to operate despite communication breakdown.

→ The CAP theorem states that it is not possible to guarantee all 3 of the desirable properties at the same time

Classification based on CAP characteristics

(i) CP database - consistency & partition tolerance at the cost of availability.

→ If partition occurs, system shuts down the non-consistent node until the partition is resolved.
e.g. BigTable

MongoDB

Berkeley DB

(ii) AP database

→ If partition occurs, nodes remain available, but those at the wrong end of a partition might get an older version
e.g. Voldemort, Kai

(iii) CA database - If there is partition, cannot deliver fault tolerance

e.g. Vertica

Greenplum

Document-based Database Systems - MongoDB

- A MongoDB database may have one or more collections
 - group of MongoDB documents
equiv to a RDBMS's table
all documents in a collection
are related.
- A collection may have zero or more documents
 - a set of key value pairs
have a dynamic schema
equivalent to a row tuple in an RDBMS
- need not have the same fields
- documents are addressed by a unique key → provided by
 - MongoDB itself
 - 16-bit field
 - automatically indexed
- Documents are stored in JSON format

Functions

db.createCollection

Insertion: db.mycoll.insert()

db.<collection name>.remove (<condn>)
say: \$title: }

Characteristics

- (i) uses 2-phase commit protocol
- (ii) replication - master-slave approach
arbiter
- (iii) Sharding
- (iv) Accessing: hashing / range partitioning